

# EC50 Prelim Results

March 26, 2017

## 1 OSM COMPETITION: A predictive model for EC50 Whole Cell Potency.

The aim of the competition is to develop a computational model that predicts which molecules will block the malaria parasite's ion pump, PfATP4.

Submitted by James McCulloch - james.duncan.mcculloch@gmail.com

### 1.1 Preliminary Results

### 1.2 Using EC50 Whole Cell Potency $\leq 500\text{nMol}$ to predict ION\_ACTIVITY gives a test set AUC score of 0.95 using a custom Neural Network and AUC score of 0.93 using an "off-the-shelf" logistic classifier; ION\_ACTIVITY[ACTIVE, INACTIVE].

See the discussion at the bottom of the notebook - just above the similarity maps.

1.2.1 The test set prediction that whole cell potency EC50  $\leq 1000\text{nMol}$  AUC score was 0.82; [Active, Inactive].

1.2.2 The test set prediction that whole cell potency EC50  $\leq 500\text{nMol}$  AUC score was 0.93; [Active, Inactive].

1.2.3 The test set prediction that whole cell potency EC50  $\leq 200\text{nMol}$  AUC score was 0.82; [Active, Inactive].

The ML classifiers use molecular Morgan RDKit fingerprints.

The results appear robust as they are duplicated using an "off-the-shelf" SKLearn classifier and a custom Neural Net classifier.

**Data Preliminaries.** The dataset was prepared from the supplied spreadsheet "ION Regulation Data for OSM Competition.xls". No additional compounds (e.g. from the Master Spreadsheet) were used.

Records were to be separated into "TRAIN" or "TEST" records for classification. All records with an "Ion Regulation Test Set" column of "B" or "C" were tagged as "TEST".

## 2 1. EC50 Whole Cell Potency

Data for EC50 analysis.

The EC50 value from the column "Potency vs Parasite (uMol)" was used. If this column is blank then the assays from the column "Alternative EC50 from ChEMBL (uM)" were used. If there was more than one assay in the column "Alternative EC50 from ChEMBL (uM)" (separated by semi-colons), then the average of these was used. All rows where there was not an available EC50 value were left blank (and will be filtered at runtime). Any "Potency Qualifier" modifiers were ignored, these are irrelevant for classifiers as all were for EC50 values of  $\geq 10$  uMol (all "INACTIVE").

### 2.0.4 EC50 $\leq 1000$ nMol Analysis

The EC50\_1000 field was generated by classifying all records with EC50 values  $\leq 1$  uMol as "ACTIVE" else "INACTIVE".

### 2.0.5 EC50 $\leq 500$ nMol Analysis

The EC50\_500 field was generated by classifying all records with EC50 values  $\leq 0.5$  uMol as "ACTIVE" else "INACTIVE".

### 2.0.6 EC50 $\leq 200$ nMol Analysis

The EC50\_200 field was generated by classifying all records with EC50 values  $\leq 0.2$  uMol as "ACTIVE" else "INACTIVE".

From a total of 703 valid records in the competition spreadsheet (SMILES available) there were 646 "TRAIN" records and 37 "TEST" records with a valid EC50 column.

**EC50 analysis methodology.** A DNN [2048, 2048, 512, 64, 3] and SKlearn logistic classifier (LOGC) were trained against Morgan finger prints (mol radius = n) (tabulated below) of the "TRAIN" records then the 37 "TEST" records were presented to the trained ML classifiers. The results of the ability of the ML classifier to determine the ION\_ACTIVITY status of "TEST" molecules are tabulated below as a Area Under Curve (AUC) statistic.

### 2.0.7 Results Summary

```
In [1]: from IPython.display import display
import pandas as pd
print("EC50_1000 Results")
ec50_1000_results = pd.read_csv("./EC50_1000_results.csv")
display(ec50_1000_results)
print("EC50_500 Results")
ec50_500_results = pd.read_csv("./EC50_500_results.csv")
display(ec50_500_results)
print("EC50_200 Results")
ec50_200_results = pd.read_csv("./EC50_200_results.csv")
display(ec50_200_results)
```

EC50\_1000 Results

	MODEL	M2048_1	M2048_2	M2048_3	M2048_4	M2048_5	DRAGON
0	LOGC	0.70	0.7	0.73	0.79	0.81	0.61
1	DNN	0.61	0.7	0.73	0.81	0.82	0.63

#### EC50\_500 Results

	MODEL	M2048_1	M2048_2	M2048_3	M2048_4	M2048_5	M2048_6	DRAGON
0	LOGC	0.78	0.81	0.86	0.84	0.91	0.86	0.59
1	DNN	0.74	0.82	0.84	0.88	0.93	0.88	0.72

#### EC50\_200 Results

	MODEL	M2048_1	M2048_2	M2048_3	M2048_4	M2048_5	DRAGON
0	LOGC	0.67	0.71	0.70	0.69	0.70	0.63
1	DNN	0.77	0.82	0.76	0.68	0.69	0.80

Where the ML MODELS are as follows:

1. LOGC - A Logistic classifier from SKLearn.
2. DNN - A Deep Neural Network classifier [2048, 2048, 512, 64, 3] from the Keras toolkit. Cross-entropy loss function.

The molecular descriptors are as follows:

1. M2048\_n - A 2048 bit, mol radius = n, Morgan fingerprint from RDKit
2. DRAGON - A vector of 1666 pharmacophore molecular fields scaled to be on the interval [0,1].

### 2.0.8 Modelling.

To run these models, download (follow the readme setup) the software on gitHub here: <https://github.com/kellerberrin/OSM-QSAR>

For the LOGC SKLearn model (-help for flag descriptions) the following cmd was used:

```
$python OSM_QSAR.py --classify logc --depend EC50_500 --indep MORGAN2048_5
```

For the DNN with MORGAN fingerprints:

```
$python OSM_QSAR.py --classify bin_m --depend EC50_500 --indep MORGAN2048_5 --train
```

For the DNN with the DRAGON descriptor:

```
$python OSM_QSAR.py --classify bin_d --depend EC50_500 --indep DRAGON --train 300 -
```

Convergence of the DNN was prompt in all cases and occurred within 300 epochs.

## Commentary.

- Morgan fingerprints out-perform DRAGON in all cases.
- LOGC and DNN have very similar performance with Morgan fingerprints.
- A big jump in AUC from 0.82 training on EC50\_200 (active at 200nM) to 0.93 training on EC50\_500 (active at 500nM)
- The LOGC classifier got significantly lower AUC on DRAGON than the DNN. This reinforces our observation in the previous notebook that the “off-the-shelf” classifiers of SKLearn have trouble interpreting the unstructured data of DRAGON.

We now take a closer look at the results.

The following fields are available in the results files.

- ID
- EC50 - in uMol
- ION\_ACTIVITY - molecule PfATP4 ion activity [ACTIVE, INACTIVE, PARTIAL, X\_MISSING] (X\_MISSING is unknown ion activity)
- ACTUAL\_200 - [ACTIVE, INACTIVE] at 200 nMol
- ACTUAL\_500 - [ACTIVE, INACTIVE] at 500 nMol
- ACTUAL\_1000 - [ACTIVE, INACTIVE] at 1000 nMol
- M2\_200\_90 - The probability map of the DNN trained 90 epochs with –depend EC50\_200 –indep MORGAN2048\_2 (AUC 0.82)
- M5\_500\_40 - The probability map of the DNN trained 40 epochs with –depend EC50\_500 –indep MORGAN2048\_5 (AUC 0.87)
- M5\_500\_130 - The probability map of the DNN trained 130 epochs with –depend EC50\_500 –indep MORGAN2048\_5 (AUC 0.90)
- M5\_500\_250 - The probability map of the DNN trained 250 epochs with –depend EC50\_500 –indep MORGAN2048\_5 (AUC 0.93)
- M5\_1000\_210 - The probability map of the DNN trained 210 epochs with –depend EC50\_1000 –indep MORGAN2048\_5 (AUC 0.81)
- LOGC\_500 - The probability map of the SKLearn logistic classifier LOGC trained against EC50\_500 (AUC 0.91)
- SMILE

Note that M5\_500\_40, M5\_500\_130 and M5\_500\_250 are all from the same DNN training session.

First, as a sanity check, we examine the Pearson correlations of the EC50\_500 probability maps.

```
In [81]: # Import the probability maps
import pandas as pd
test_results = pd.read_csv("./test_results.csv")
train_results = pd.read_csv("./train_results.csv")
EC50_500_active = ["M5_500_40", "M5_500_130", "M5_500_250", "LOGC_500"]
all_active = EC50_500_active + ["M5_1000_210", "M2_200_90"]
```

```
In [47]: # Train results Pearson correlation
corr = train_results[EC50_500_active].corr(method="pearson")
corr
```

```
Out [47]:
```

	M5_500_40	M5_500_130	M5_500_250	LOGC_500
M5_500_40	1.000000	0.999100	0.979215	0.988947
M5_500_130	0.999100	1.000000	0.980659	0.989751
M5_500_250	0.979215	0.980659	1.000000	0.970623
LOGC_500	0.988947	0.989751	0.970623	1.000000

```
In [46]: # Test results Pearson correlation
corr = test_results[EC50_500_active].corr(method="pearson")
corr
```

```
Out [46]:
```

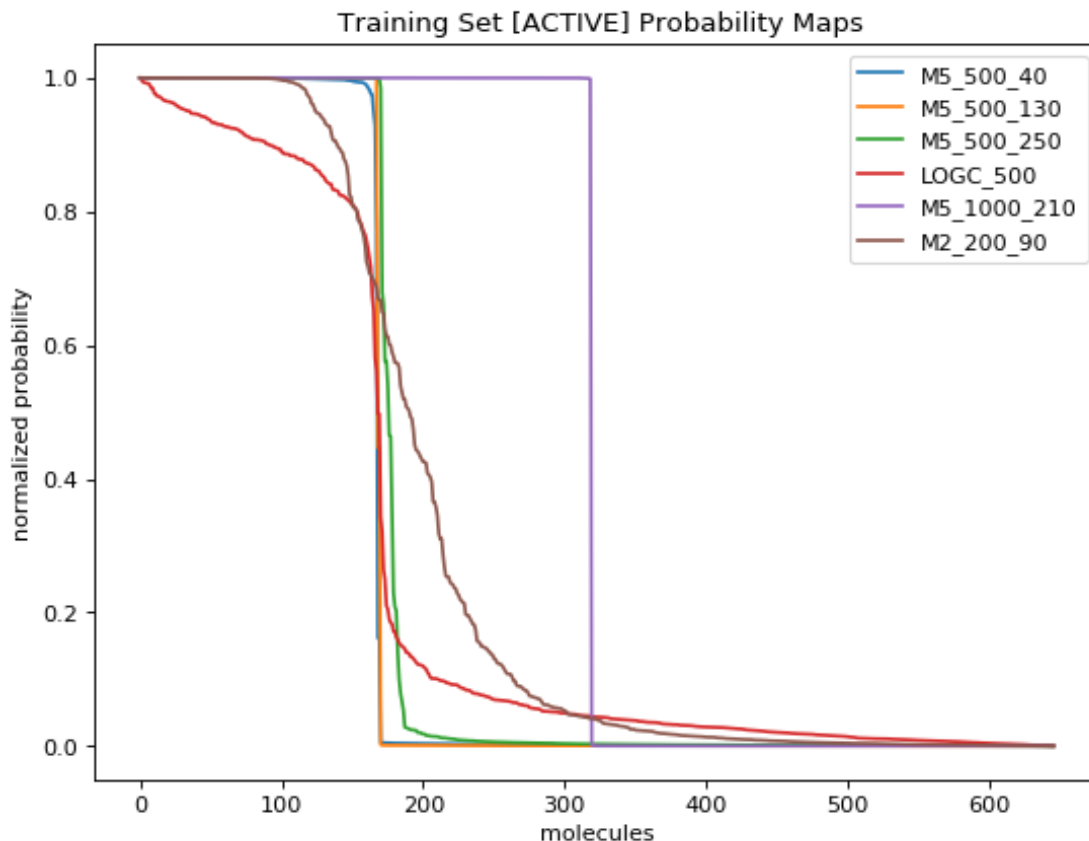
	M5_500_40	M5_500_130	M5_500_250	LOGC_500
M5_500_40	1.000000	0.736947	0.543742	0.706448
M5_500_130	0.736947	1.000000	0.784117	0.849912
M5_500_250	0.543742	0.784117	1.000000	0.785143
LOGC_500	0.706448	0.849912	0.785143	1.000000

We visualize the training set probability maps by normalizing them to the unit interval [0, 1] and sorting them in descending order.

```
In [48]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from pylab import *
from sklearn.preprocessing import minmax_scale

def sort_map(column):
    array = minmax_scale(train_results[column])
    return array[np.argsort(-array)]

scale = 1.0
fig = plt.figure(num=None, figsize=(8 * scale, 6 * scale), dpi=80, facecol
for map in all_active: plt.plot(sort_map(map), label=map)
xlabel("molecules")
ylabel("normalized probability")
title(" Training Set [ACTIVE] Probability Maps")
legend(loc=1); # upper right corner
```



### Commentary.

- Very rapid DNN convergence. We see that the probability map for M5\_500\_40 has converged to a step function in 40 epochs.
- All of the probability maps with the exception of M2\_200\_90 approximate step functions.
- **Key question**, What exemplar molecules were added to the [ACTIVE] training set when moving from EC50\_200 to EC50\_500?

## 2.1 Molecular Classification.

This section examines training and test molecules that are classified by EC50\_200, EC50\_500 and EC50\_1000. The python code comes from Greg Landrum's blog posts. You must be in an anaconda (or other) python environment that has previously setup rdkit to execute this code.

The python code below generates the following data\_frames:

- ec50\_200\_active [ACTIVE] training molecules in EC50\_200
- ec50\_500\_active [ACTIVE] training molecules in EC50\_500
- ec50\_500\_added the added [ACTIVE] training molecules ec50\_500\_active - ec50\_200\_active
- ec50\_1000\_active [ACTIVE] training molecules in EC50\_1000

- ec50\_1000\_added the added [ACTIVE] training molecules ec50\_1000\_active - ec50\_500\_active

The code also generates the rdkit mols and labels for use is displaying moles.

The labels have the following form: "1 MMV669304 A A 280 (1.000)"

- The number "1" is the position of the molecule in the data\_frame. This is only meaningful if the data\_frame has been sorted.
- Molecular ID "MMV669304"
- The first "A" indicates if the molecule is [ACTIVE] at EC50 <= 500nMol. Values are [ACTIVE] = "A", [INACTIVE] = "I".
- The second "A" is the PfATP4 ION\_ACTIVITY class. Values are [ACTIVE] = "A", [INACTIVE] = "I", [PARTIAL] = "P", no ION\_ACTIVITY assay = "X".
- "280" is the EC50 assay in nMols.
- "(1.000)" is the M5\_500\_250 probability map value of the molecule. For training set molecules this will always be close 1.000.

```
In [98]: ec50_200_active = train_results.loc[train_results["ACTUAL_200"] == "ACTIVE"]
ec50_200_active_list = ec50_200_active["ID"].tolist()
ec50_500_active = train_results.loc[train_results["ACTUAL_500"] == "ACTIVE"]
ec50_500_active_list = ec50_500_active["ID"].tolist()
added_ec50_500_list = list(set(ec50_500_active_list) - set(ec50_200_active_list))
ec50_500_added = train_results.loc[train_results["ID"].isin(added_ec50_500_list)]
ec50_1000_active = train_results.loc[train_results["ACTUAL_1000"] == "ACTIVE"]
ec50_1000_active_list = ec50_1000_active["ID"].tolist()
added_ec50_1000_list = list(set(ec50_1000_active_list) - set(ec50_500_active_list))
ec50_1000_added = train_results.loc[train_results["ID"].isin(added_ec50_1000_list)]

def mol_label_list(data_frame): # Function to produce rdkit mols and associated labels
    id = data_frame["ID"].tolist()
    klass = data_frame["ACTUAL_500"].tolist()
    potency = data_frame["EC50"].tolist()
    ion_activity = data_frame["ION_ACTIVITY"].tolist()
    map_prob = data_frame["M5_500_250"].tolist()
    labels = []
    for idx in range(len(id)):
        labels.append("{} {} {} {} {:.5.0f} {:.5.4f}".format(idx+1, id[idx], klass[idx][0],
                                                                ion_activity[idx],
                                                                potency[idx]*100))

    smiles = data_frame["SMILE"].tolist()
    mols = [Chem.MolFromSmiles(smile) for smile in smiles]
    return mols, labels

print("len(ec50_200_active):", len(ec50_200_active), "len(ec50_500_active):", len(ec50_500_active),
      "len(ec50_500_added):", len(ec50_500_added), "len(ec50_1000_active):", len(ec50_1000_active),
      "len(ec50_1000_added):", len(ec50_1000_added))
```

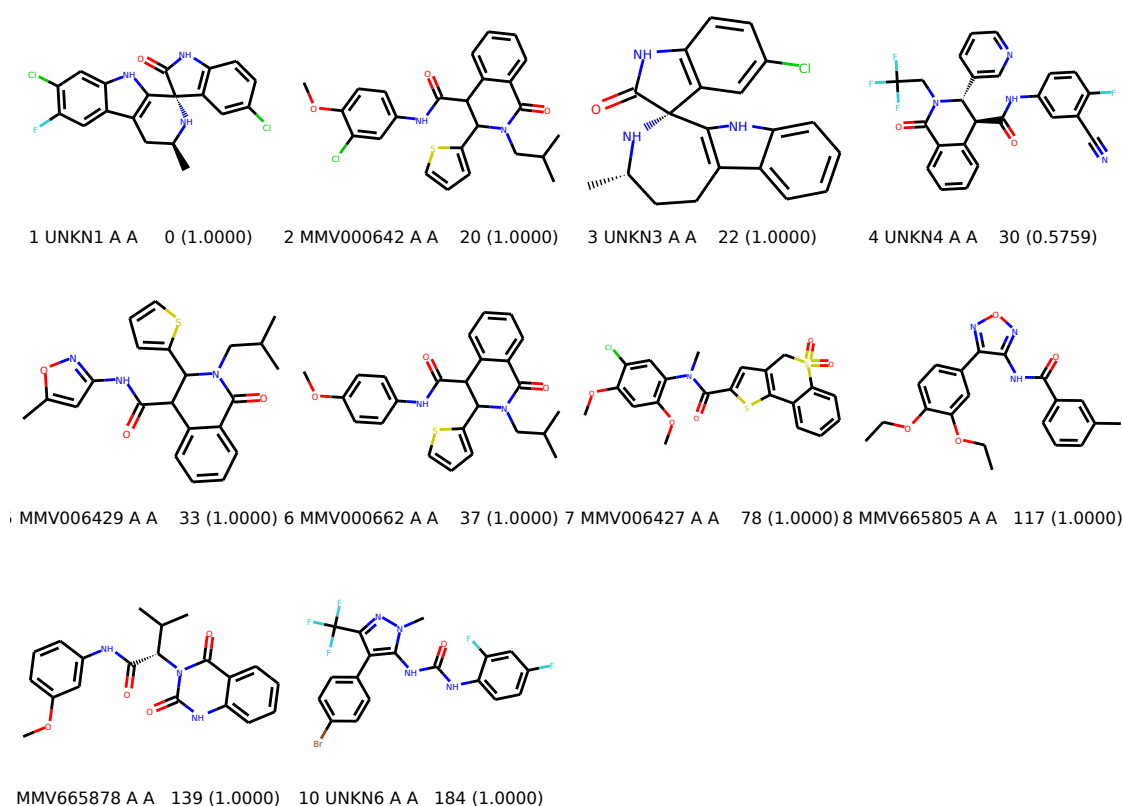
```
len(ec50_200_active): 71 len(ec50_500_active): 169 len(ec50_500_added): 98 len(ec50_1000_active): 169 len(ec50_1000_added): 98
```

```
In [84]: from rdkit import Chem
         from rdkit.Chem import Draw
         from rdkit.Chem.Draw import IPythonConsole
         from rdkit import rdBase
         IPythonConsole.ipynon_useSVG=True
```

## ION\_ACTIVITY [ACTIVE] in EC50\_200

```
In [99]: ion_active = ec50_200_active.loc[train_results["ION_ACTIVITY"] == "ACTIVE"]
         mols, labels = mol_label_list(ion_active)
         Draw.MolsToGridImage(mols, legends=labels, molsPerRow=4)
```

Out [99]:




## ION\_ACTIVITY [ACTIVE] Exemplar molecules that were added to the training set when moving from EC50\_200 to EC50\_500

```
In [100]: ion_active = ec50_500_added.loc[train_results["ION_ACTIVITY"] == "ACTIVE"]
          mols, labels = mol_label_list(ion_active)
          Draw.MolsToGridImage(mols, legends=labels, molsPerRow=4)
```

Out [100]:






**Commentary** These molecules have the same Triazole arm as we noticed in the previous notebook when trying to classify the molecular ION\_ACTIVITY using D840\_ACTIVE (DRAGON). This structure is also well represented in the test molecules.

**ION\_ACTIVITY [ACTIVE] Exemplar molecules that were added to the training set when moving from EC50\_500 to EC50\_1000**

```
In [101]: ion_active = ec50_1000_added.loc[train_results["ION_ACTIVITY"] == "ACTIVE"]
          mols, labels = mol_label_list(ion_active)
          Draw.MolsToGridImage(mols, legends=labels, molsPerRow=4)
```


Out [101]:



**The results of the EC50\_500 classification of the test molecules.**

```
In [104]: sorted = test_results.sort_values("M5_500_250", ascending=False)
          mols, labels = mol_label_list(sorted)
          Draw.MolsToGridImage(mols, legends=labels, molsPerRow=4)
```

Out[104]:



## Commentary

- The immediate and possibly very important observation is the “false positive” test molecules that were incorrectly identified as ION\_ACTIVITY [ACTIVE] by the D840\_ACTIVE (DRAGON) classifier in the previous notebook; OSM-S-367 and OSM-S-201 [INACTIVE] and OSM-S-204 [PARTIAL] are now classified as very low probability of being EC50\_500 [ACTIVE] by the M5\_500\_250 classifier.
- This may be the link we are looking for between ION\_ACTIVITY and EC50. And suggests a two stage selection methodology.
  1. Use the D840\_ACTIVE (DRAGON) classifier described in the previous notebook to select likely molecular structures, including false positive candidates such as OSM-S-367, OSM-S-201 and OSM-S-204.
  2. Then triage the candidates by predicting EC50\_500 activity using M5\_500\_250. This would then exclude OSM-S-367, OSM-S-201 and OSM-S-204.
- What about “false negatives”. OSM-S-371 and OSM-S-370 are ION\_ACTIVITY [ACTIVE] but are classified by D840\_ACTIVE as 31/35 and 32/35 ([INACTIVE]) respectively because the structure of the Pyridine arm of these molecules is quite different to that of highly ranked molecules. However, M5\_500\_250 classifies these molecules as EC50\_500 [ACTIVE] (prob 1.000). Using the procedure described above these molecules would be excluded at the first round. Although this represents wasted opportunity, not wasted effort.
- We could just use the EC50\_500 molecule classification as a proxy for ION\_ACTIVITY classification. This would solve the “false negative” problem but introduces new “false positives” such as OSM-S-366.

**Using EC50\_500 to predict ION\_ACTIVITY** The python code below uses the molecule ION\_ACTIVITY class and the probability map from M5\_500\_250 and LOGC\_500 to calculate the AUC score we would generate by using EC50\_500 to predict ION\_ACTIVITY.

```
In [112]: from sklearn.metrics import roc_auc_score

# Get the M5_500_250 prob vector

prob = test_results["M5_500_250"].tolist()

# Get the LOGC_500 prob vector

logc_prob = test_results["LOGC_500"].tolist()

# Get the ION ACTIVITY vector

ion = test_results["ION_ACTIVITY"].tolist()
```

```

# Filter out molecules without an ion assay and convert ION_ACTIVITY to a

filter_prob = []
filter_logc_prob = []
filter_ion = []
for idx in range(len(prob)):

    if ion[idx] == "ACTIVE":
        filter_prob.append(prob[idx])
        filter_logc_prob.append(logc_prob[idx])
        filter_ion.append(1)
    elif ion[idx] == "INACTIVE" or ion[idx] == "PARTIAL":
        filter_prob.append(prob[idx])
        filter_logc_prob.append(logc_prob[idx])
        filter_ion.append(0)

ion_auc = roc_auc_score(filter_ion, filter_prob, average=None, sample_weight=None)

ion_logc_auc = roc_auc_score(filter_ion, filter_logc_prob, average=None, sample_weight=None)

print("AUC when using M5_500_250 to predict ION_ACTIVITY:", ion_auc)
print("AUC when using LOGC_500 to predict ION_ACTIVITY:", ion_logc_auc)

```

AUC when using M5\_500\_250 to predict ION\_ACTIVITY: 0.947712418301  
AUC when using LOGC\_500 to predict ION\_ACTIVITY: 0.928104575163

## Commentary

### 2.2 Using EC50\_500 to predict ION\_ACTIVITY generates an AUC = 0.95 using a custom Neural Network and AUC = 0.93 using an “off-the-self” logistic classifier.

This could be the solution to the problem of classifying ION\_ACTIVITY.


**Similarity Maps** The LOGC\_500 similarity maps (step functions such as M5\_500\_250 may not work well with similarity maps) for OSM-S-379 [ACTIVE], OSM-S-367 [INACTIVE], OSM-S-201 [INACTIVE], OSM-S-204[INACTIVE], OSM-S-370 [ACTIVE] and OSM-S-371 [ACTIVE] are displayed below.

OSM-S-367


OSM-S-379

OSM-S-201


OSM-S-204



OSM\_S\_367




OSM\_S\_379

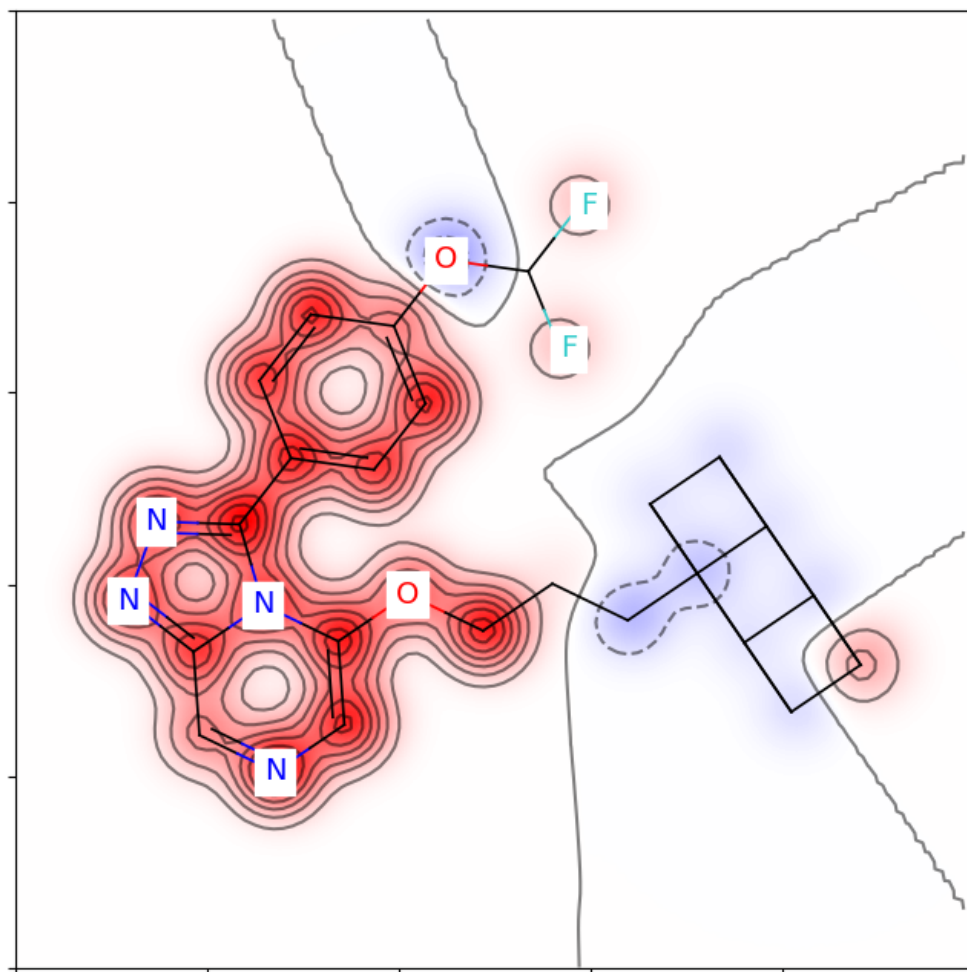


OSM\_S\_201



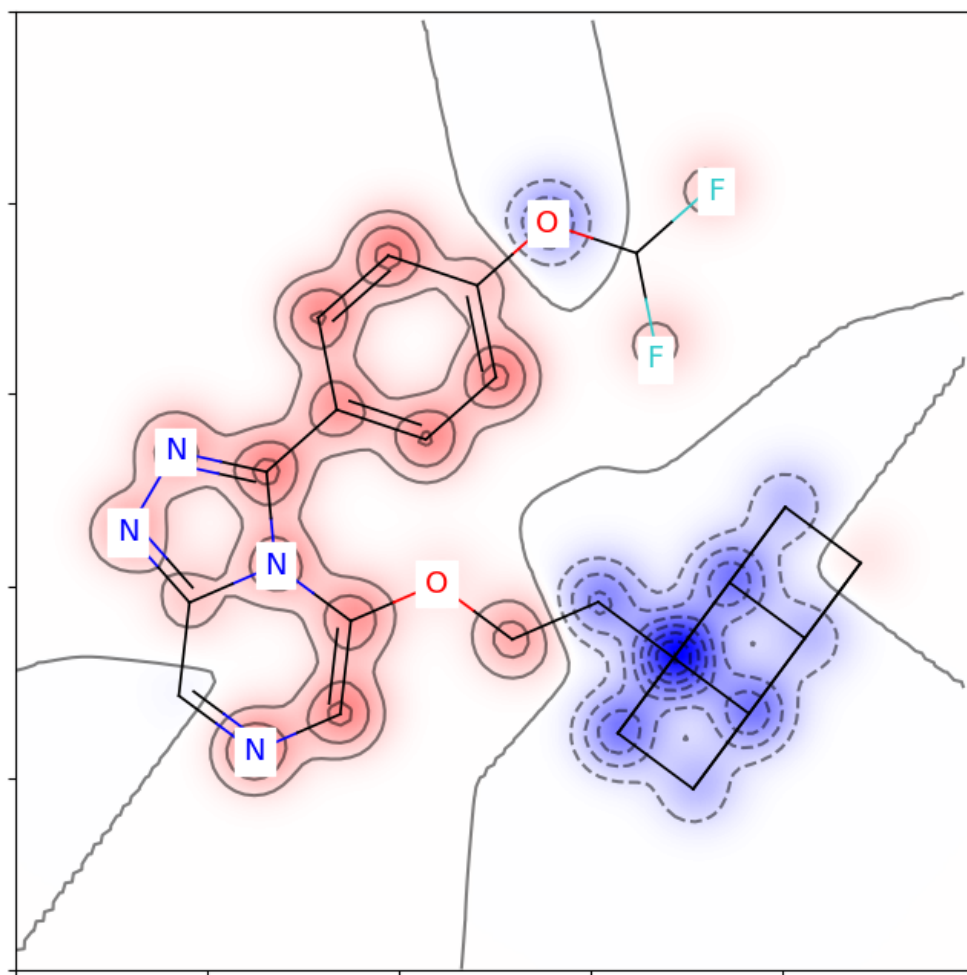


OSM\_S\_204



OSM\_S\_370

OSM-S-370



OSM\_S\_371

OSM-S-371